

Parallel Computing

Degree in Computer Science Engineering (ETSIINF)

Year 2024-25 ◇ Partial exam 27/1/25 ◇ Block OpenMP ◇ Duration: 1h 30m



UNIVERSITAT
POLITÀCNICA
DE VALÈNCIA

Question 1 (1 point)

Given the following code:

```
double function( int n, double v[], double w[], double *a ) {
    int i, j, c = 0; double b, f = 0, e;
    for (i=0; i<n; i++) {
        e = 0;
        for (j=i+1; j<n; j++) {
            b = sqrt(v[i]*w[j]);
            if ( b > e ) e = b;
            c++;
        }
        v[i] = e;
        if ( e > f ) f = e;
    }
    *a = f;
    return c;
}
```

0.3 p.

- (a) Parallelize the outermost loop.

Solution: We just insert the following directive right before the first loop:

```
#pragma omp parallel for private(e,j,b) reduction(max:f) reduction(+:c)
```

0.3 p.

- (b) Parallelize the innermost loop.

Solution: We just insert the following directive right before the second loop:

```
#pragma omp parallel for private(b) reduction(max:e) reduction(+:c)
```

0.4 p.

- (c) Compute the computational cost (in flops) of the original sequential version, assuming that the `sqrt` function has a cost of 3 Flops. Taking into account the cost of a single iteration of the `i` loop, discuss whether there would be a good load balancing in case of using `schedule(static)` in section (a). Repeat the discussion in case the same scheduling was used in section (b).

Solution:

Sequential cost:

$$t(n) = \sum_{i=0}^{n-1} \sum_{j=i+1}^{n-1} 4 \approx \sum_{i=0}^{n-1} (4n - 4i) \approx 4n^2 - 4\frac{n^2}{2} = 2n^2 \text{ flops.}$$

The cost of an iteration of the `i` loop is approximately $4(n - i)$. Since it depends on i , the cost of each iteration is different; in particular, the first iterations are the most costly and the last ones are the least costly. Using `schedule(static)` gives place to a load imbalance, since thread 0 will be assigned the most costly block of iterations, while the last thread will be assigned the least costly block of iterations.

In the case of loop `j`, the cost of one iterations is 4 flops. Since all iterations cost the same, `schedule(static)` will not produce load imbalance.

Question 2 (1.4 points)


Given the following program, where function **generate** modifies its argument and has a computational cost of $6N^2$ Flops.

```
float fprocess(float A[N][N], float factor) {
    int i,j; float result=0.0;
    for (i=0;i<N;i++)
        for (j=0;j<N;j++)
            result+=A[i][j]/(factor+j);
    return result;
}

void processing() {
    float A[N][N], B[N][N], n1, n2, n3, n4;
    generate(A);          // T1
    n1=fprocess(A,1.1);    // T2
    n2=fprocess(A,1.2);    // T3
    generate(B);          // T4
    n3=fprocess(B,1.1);    // T5
    n4=fprocess(B,1.2);    // T6
    printf("Result: %f\n", n1+n2+n3+n4); // T7
}
```

0.4 p.

- (a) Obtain the task dependency graph of function **processing**, calculating the critical path and its length, as well as the average and maximum degree of concurrency.

Solution:

The critical path would be $T_1 \rightarrow T_2 \rightarrow T_7$. Function **fprocess** has a cost of $3N^2$ Flops, and therefore the critical path has a length of $9N^2 + 3$. The maximum degree of concurrency is 4 and the average degree of concurrency is $M = \frac{24N^2+3}{9N^2+3} \approx 2.67$

0.6 p.

- (b) Write an efficient parallel implementation of function **processing** based on sections.

Solution:

```
void processing() {
    float A[N][N], B[N][N];
    float n1, n2, n3, n4;

    #pragma omp parallel
    {
        #pragma omp sections
        {
            #pragma omp section
            generate(A);

```

```

        #pragma omp section
        generate(B);
    }
    #pragma omp sections
    {
        #pragma omp section
        n1=fprocess(A,1.1);
        #pragma omp section
        n2=fprocess(A,1.2);
        #pragma omp section
        n3=fprocess(B,1.1);
        #pragma omp section
        n4=fprocess(B,1.2);
    }
}
printf("Result: %f\n", n1+n2+n3+n4);
}

```

0.4 p.

- (c) Compute the sequential computational cost. Compute the parallel cost, the speed-up and the efficiency in the case of using 2 threads, and also in the case of using 4 threads.

Solution:

$$t(N) = 24N^2 \text{ Flops}$$

$$t(N, 2) = 6N^2 + 3N^2 + 3N^2 + 3 \approx 12N^2 \text{ Flops}$$

$$Sp(N, 2) = \frac{24N^2}{12N^2} = 2$$

$$E(N, 2) = \frac{2}{2} = 1$$

$$t(N, 4) = 6N^2 + 3N^2 + 3 \approx 9N^2 \text{ Flops}$$

$$Sp(N, 4) = \frac{24N^2}{9N^2} = 2,67$$

$$E(N, 4) = \frac{2,66}{4} = 0,67$$

Question 3 (1.1 points)

The following program obtains an approximation of the PI number from the generation of 200000 points, calculating the number of points that would lie inside or outside a circumference of radius r . From that ratio we obtain an approximation of the area and, therefore, of PI. Function `random(0,N)` returns an integer random number between 0 and $N - 1$.

```

#define N 20
#define SAMPLES 200000
int main() {
    int i,j, ipos, jpos, imax, jmax, imin, jmin, r, min, max;
    int A[N][N], inside=0, outside=0;

    for (i=0;i<N;i++)
        for (j=0;j<N;j++)
            A[i][j]= 0;
}

```

```

for (i=0;i<SAMPLES;i++) {
    ipos = random(0,N);
    jpos = random(0,N);
    A[ipos][jpos]++;
}
min = A[0][0]; max = A[0][0]; r = N/2;
for (i=0;i<N;i++) {
    for (j=0;j<N;j++) {
        if ((r-i)*(r-i)+(r-j)*(r-j)<r*r)
            inside+=A[i][j];
        else
            outside+=A[i][j];

        if (min>A[i][j]) {
            min=A[i][j];
            imin = i;
            jmin = j;
        }
        if (max<A[i][j]) {
            max=A[i][j];
            imax = i;
            jmax = j;
        }
    }
}
printf("d=%d, f=%d\n", inside, outside);
printf("Pi = %f\n", (4.0*inside)/(inside+outside));
printf("Max A[%d][%d] = %d \n", imax,jmax,max);
printf("Min A[%d][%d] = %d \n", imin,jmin,min);
return 0;
}

```

0.7 p.

- (a) Write a parallel implementation using OpenMP. It is not necessary to parallelize the loop that initializes matrix A. It is allowed to use several parallel regions.

Solution:

```

#define N 20
#define SAMPLES 200000
int main() {
    int i,j, ipos, jpos, imax, jmax, imin, jmin, r;
    int min, max;
    int A[N][N];
    int inside=0, outside=0;

    for (i=0;i<N;i++)
        for (j=0;j<N;j++)
            A[i][j] = 0;

    #pragma omp parallel for private (ipos,jpos)
    for (i=0;i<SAMPLES;i++) {

```

```

        ipos = random(0,N);
        jpos = random(0,N);
        #pragma omp atomic
        A[ipos][jpos]++;
    }
    min = A[0][0];
    max = A[0][0];
    r = N/2;
    #pragma omp parallel for private (j) reduction (+:inside, outside)
    for (i=0;i<N;i++) {
        for (j=0;j<N;j++) {
            if ((r-i)*(r-i)+(r-j)*(r-j)<r*r)
                inside+=A[i][j];
            else
                outside+=A[i][j];

            if (min>A[i][j])
                #pragma omp critical (min)
                if (min>A[i][j]) {
                    min=A[i][j];
                    imin = i;
                    jmin = j;
                }
            if (max<A[i][j])
                #pragma omp critical (max)
                if (max<A[i][j]) {
                    max=A[i][j];
                    imax = i;
                    jmax = j;
                }
        }
    }
    printf("d=%d, f=%d\n", inside, outside);
    printf("Pi = %f\n", (4.0*inside)/(inside+outside));
    printf("Max A[%d][%d] = %d \n", imax,jmax,max);
    printf("Min A[%d][%d] = %d \n", imin,jmin,min);
    return 0;
}

```

0.4 p.

- (b) Modify the program so that it prints the value of variables `inside` and `outside` that have been computed by each of the threads.

Solution:

```

...
r = N/2;
/* Up to this point, same code as in section a */
#pragma omp parallel private (j) reduction (+:inside, outside)
{
    #pragma omp for
    for (i=0;i<N;i++) {
        /* Inside the loop, same code as in section a */

```

```
        ...
    }
    printf("Values of inside and outside for thread %d: %d and %d\n",
           omp_get_thread_num(), inside, outside);
}
/* From here, same code as the original program */
printf("d=%d, f=%d\n", inside, outside);
...
```