

0.3 p.

- (c) Calculate the sequential time, parallel time, speedup and efficiency corresponding to a single iteration of loop k , assuming that only the first inner loop has been parallelized (first loop with i).

Solution:

$$t(N) = \sum_{i=0}^{N-1} \sum_{j=0}^{N-1} 2 + 1 + \sum_{i=0}^{N-1} \sum_{j=0}^{i-1} 2 \approx \sum_{i=0}^{N-1} 2N + 2 \sum_{i=0}^{N-1} i \approx 2N^2 + 2 \frac{N^2}{2} = 3N^2 \text{ flops.}$$

$$t(N, p) = \sum_{i=0}^{\frac{N}{p}-1} \sum_{j=0}^{N-1} 2 + \sum_{i=0}^{N-1} \sum_{j=0}^{i-1} 2 \approx \sum_{i=0}^{\frac{N}{p}-1} 2N + N^2 = \frac{2N^2}{p} + N^2 = \left(\frac{2}{p} + 1 \right) N^2 \text{ flops.}$$

$$S(N, p) = \frac{t(N)}{t(N, p)} = \frac{3N^2}{\left(\frac{2}{p} + 1 \right) N^2} = \frac{3}{\frac{2}{p} + 1}.$$

$$E(N, p) = \frac{S(N, p)}{p} = \frac{3}{2 + p}.$$


Question 2 (1.2 points)

Given the following function, where each of the invoked functions modifies only the vector that it receives as first argument:

```
void func(double a[], double b[], double c[], double d[], int n) {
    double x;
    task1(b,a,n);      /* Cost: 4n flops */
    x=task2(c,a,n);    /* Cost: 3n flops */
    task3(b,c,n);      /* Cost: 2n flops */
    task4(d,a,n);      /* Cost:  n flops */
    task5(d,a,n);      /* Cost:  n flops */
    task6(d,x,n);      /* Cost: 3n flops */
}
```

0.4 p.

- (a) Draw the task dependency graph. Mark the critical path indicating its length and obtain the average degree of concurrency.

Solution:

Critical path: $T1 \rightarrow T3$ or also $T2 \rightarrow T6$.

Length of critical path: $L = 4n + 2n = 6n$ flops

Average degree of concurrency:

$$M = \frac{4n + 3n + 2n + n + n + 3n}{6n} = \frac{14n}{6n} = 2.33$$

0.6 p.

- (b) Write a parallel version using OpenMP. The execution time must be minimized.

Solution:

```

void func(double a[], double b[], double c[], double d[], int n) {
    double x;
    #pragma omp parallel
    {
        #pragma omp sections
        {
            #pragma omp section
            task1(b,a,n);
            #pragma omp section
            x=task2(c,a,n);
            #pragma omp section
            {
                task4(d,a,n);
                task5(d,a,n);
            }
        }
        #pragma omp sections
        {
            #pragma omp section
            task3(b,c,n);
            #pragma omp section
            task6(d,x,n);
        }
    }
}

```

0.2 p.

- (c) Obtain the speedup of the parallel version when run with 3 processors.

Solution: Sequential execution time:

$$t(n) = 4n + 3n + 2n + n + n + 3n = 14n \text{ flops}$$

The parallel execution time is $4n$ flops for the first **sections**, plus $3n$ for the second **sections**:

$$t(n,p) = 4n + 3n = 7n \text{ flops}$$

Speedup:

$$S(n,p) = \frac{14n}{7n} = 2$$

Question 3 (1.2 points)

Function `pluviometry` computes different statistical values related to water precipitation occurred during a set of `ND` days in a territory for a total of `NM` months. In turn, function `month_day` provides the identifier of the month to which a certain day belongs. Parallelize the function by means of OpenMP employing just one parallel region. The ordering of days in vector `risk_days` is not relevant for the parallel version.

```

void pluviometry(float precipitation[ND],int risk_litres,float average_rain_month[NM]) {
    int i, month, nrainy_days=0, max_day, nrainy_days_month[NM];
    int nrisk_days=0, risk_days[ND];
    float sum_rain=0, precipitation_max=0, average_rain;
    float sum_rain_month[NM];

```

```

/* Initialize the vectors entries to 0 */
...

for (i=0;i<ND;i++) {
    if (precipitation[i]>0) {
        sum_rain+=precipitation[i];
        nrainy_days++;
        if (precipitation[i]>risk_litres) {
            risk_days[nrisk_days]=i+1;
            nrisk_days++;
        }
        if (precipitation[i]>precipitation_max) {
            precipitation_max=precipitation[i];
            max_day=i;
        }
        month=month_day(i);
        nrainy_days_month[month]++;
        sum_rain_month[month]+=precipitation[i];
    }
}
average_rain=sum_rain/nrainy_days;
for (i=0;i<NM;i++)
    average_rain_month[i]=sum_rain_month[i]/nrainy_days_month[i];

/* More code */
...
}

```

Solution:

```

void pluviometry(float precipitation[ND],int risk_litres,float average_rain_month[NM]) {
    int i, month, nrainy_days=0, max_day, nrainy_days_month[NM];
    int nrisk_days=0, risk_days[ND];
    float sum_rain=0, precipitation_max=0, average_rain;
    float sum_rain_month[NM] ;

    /* Initialize the vectors entries to 0 */
    ...

    #pragma omp parallel
    {
        #pragma omp for private(month) reduction(+:sum_rain,nrainy_days)
        for (i=0;i<ND;i++) {
            if (precipitation[i]>0) {
                sum_rain+=precipitation[i];
                nrainy_days++;
                if (precipitation[i]>risk_litres) {
                    #pragma omp critical (risk)
                    {

```

```

        risk_days[nrisk_days]=i+1;
        nrisk_days++;
    }
}
if (precipitation[i]>precipitation_max) {
    #pragma omp critical (maximum)
    {
        if (precipitation[i]>precipitation_max) {
            precipitation_max=precipitation[i];
            max_day=i;
        }
    }
}
month=month_day(i);
#pragma omp atomic
nrainy_days_month[month]++;
#pragma omp atomic
sum_rain_month[month]+=precipitation[i];
}
}
average_rain=sum_rain/nrainy_days;
#pragma omp for
for (i=0;i<NM;i++)
    average_rain_month[i]=sum_rain_month[i]/nrainy_days_month[i];

/* More code */
...
}

```