

**Question 1** (1.3 points)

Given the following function:

```
void fun1(double A[N][N], double x[], double y[]) {
    int i,j;

    for (i=0;i<N;i++) {
        for (j=0;j<N;j++)
            A[i][j]= x[i]*y[j];
    }
}
```

1 p.

- (a) Implement a parallel version using MPI, assuming that the input data is in process 0 and that the results must be complete in that process at the end of the execution. The problem size can be assumed to be a multiple of the number of processes.

**Solution:**

```
void fun1_par(double A[N][N], double x[N], double y[N]) {
    int p, np;
    int i,j;
    double xlcl[N];
    double Alcl[N][N];

    MPI_Comm_size(MPI_COMM_WORLD, &p);

    np = N/p;
    MPI_Scatter(x, np, MPI_DOUBLE, xlcl, np, MPI_DOUBLE, 0, MPI_COMM_WORLD);
    MPI_Bcast(y, N, MPI_DOUBLE, 0, MPI_COMM_WORLD);

    for (i=0;i<np;i++)
        for (j=0;j<N;j++)
            Alcl[i][j] = xlcl[i]*y[j];
    MPI_Gather(Alcl, np*N, MPI_DOUBLE, A, np*N, MPI_DOUBLE, 0, MPI_COMM_WORLD);
}
```

0.3 p.

- (b) Obtain the expression of the parallel execution time, indicating the communication cost of each collective operation used.

**Solution:**

$$\begin{aligned}
 t(N, p) &= t_{comm}(N, p) + t_a(N, p) \\
 t_{comm}(N, p) &= t_{Scatter} + t_{Bcast} + t_{Gather} \\
 t_{Scatter} &= (p-1)\left(t_s + \frac{N}{p}t_w\right) \\
 t_{Bcast} &= (p-1)(t_s + Nt_w) \\
 t_{Gather} &= (p-1)\left(t_s + N\frac{N}{p}t_w\right) \\
 t_a(N, p) &= \sum_{i=0}^{\frac{N}{p}-1} \sum_{j=0}^N 1
 \end{aligned}$$

$$t(N, p) = (p-1)\left(t_s + \frac{N}{p}t_w\right) + (p-1)(t_s + Nt_w) + \sum_{i=0}^{\frac{N}{p}-1} \sum_{j=0}^N 1 + (p-1)\left(t_s + N\frac{N}{p}t_w\right)$$

$$t(N, p) \approx 3pt_s + (pN + N^2)t_w + \frac{N^2}{p}$$

**Question 2** (1.1 points)

We want to send the first and last rows and columns of a rectangular matrix of size  $M \times N$  from the process identified as *root* to the rest of the processes. Below is an example for a matrix of dimensions  $M = 4$  and  $N = 5$ , where the terms identified with the symbol  $x$  correspond to all those to be sent:

$$A = \begin{pmatrix} x & x & x & x & x \\ x & \cdot & \cdot & \cdot & x \\ x & \cdot & \cdot & \cdot & x \\ x & x & x & x & x \end{pmatrix}$$

0.9 p.

- (a) Complete the body of the function whose header is included below to carry out the communication. The parameters of the function correspond to the identifier of the invoking process (*myid*), the total number of processes (*np*) and the identifier of the root process that initially has the starting *A* array and performs the communication (*root*).

```
void send_perimeter_matrix(double A[M][N], int myid, int np, int root);
```

All processes with an identifier other than *root* must store the data received in the same *A* array provided as a parameter to the function. For this purpose, point-to-point communication operations and derived data types shall be used, so as to minimize the number of transmissions to be made by the process *root* to the rest. It will be valued that no element is sent more than once to each process (especially, the elements of the 4 corners of the matrix).

**Solution:**

```
// ALTERNATIVE 1
void send_perimeter_matrix(double A[M][N], int myid, int np, int root) {
    int p;
    MPI_Datatype column;
    MPI_Datatype rows_first_last;
    MPI_Type_vector(M, 1, N, MPI_DOUBLE, &column);
    MPI_Type_vector(2, N-2, (M-1)*N, MPI_DOUBLE, &rows_first_last);
    MPI_Type_commit(&rows_first_last);
    MPI_Type_commit(&column);
    if (myid==root) {
        for (p=0; p<np; p++) {
            if (p!=root) {
                MPI_Send(A, 1, column, p, 0, MPI_COMM_WORLD);
                MPI_Send(&A[0][N-1], 1, column, p, 0, MPI_COMM_WORLD);
                MPI_Send(&A[0][1], 1, rows_first_last, p, 0, MPI_COMM_WORLD);
            }
        }
    }
    else {
        MPI_Recv(A, 1, column, root, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        MPI_Recv(&A[0][N-1], 1, column, root, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        MPI_Recv(&A[0][1], 1, rows_first_last, root, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    }
    MPI_Type_free(&rows_first_last);
    MPI_Type_free(&column);
}

// ALTERNATIVE 2
void send_perimeter_matrix(double A[M][N], int myid, int np, int root) {
    int p;
    MPI_Datatype rows_consecutive;
```

```

MPI_Datatype rows_first_last;
MPI_Type_vector(M-1,2,N,MPI_DOUBLE,&rows_consecutive);
MPI_Type_vector(2,N-1,(M-1)*N+1,MPI_DOUBLE,&rows_first_last);
MPI_Type_commit(&rows_consecutive);
MPI_Type_commit(&rows_first_last);
if (myid==root) {
    for (p=0;p<np;p++) {
        if (p!=root) {
            MPI_Send(&A[0][N-1],1,rows_consecutive,p,0,MPI_COMM_WORLD);
            MPI_Send(A,1,rows_first_last,p,0,MPI_COMM_WORLD);
        }
    }
}
else {
    MPI_Recv(&A[0][N-1],1,rows_consecutive,root,0,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
    MPI_Recv(A,1,rows_first_last,root,0,MPI_COMM_WORLD,MPI_STATUS_IGNORE);
}
MPI_Type_free(&rows_consecutive);
MPI_Type_free(&rows_first_last);
}

```

0.2 p.

(b) Obtain the communications cost.

**Solution:**

Alternative 1:

$$t_c = (p-1)(2(t_s + Mt_w) + t_s + 2(N-2)t_w)$$

Alternative 2:

$$t_c = (p-1)(t_s + 2(M-1)t_w + t_s + 2(N-1)t_w)$$

### Question 3 (1.1 points)

Given the following function, where the functions corresponding to the tasks (T1 to T6) modify only their last argument and where the cost of each of these functions is  $4N^2$  flops, except the functions T2 and T4, whose cost is  $3N^2$  flops each.

```

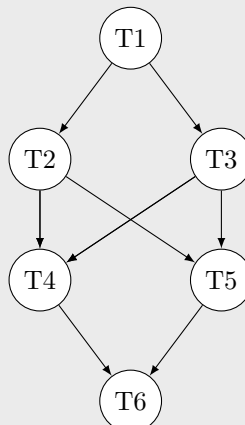
void func(double A[N][N], double w[N]) {
    double x[N],y[N],v[N],alpha;
    T1(A,x);
    T2(A,x,y);
    T3(x,v);
    T4(y,v,w);
    T5(A,y,v,&alpha);
    T6(alpha,w);
}

```

0.3 p.

(a) Draw the graph of data dependencies between tasks.

**Solution:**



0.6 p.

- (b) Implement a parallel version with MPI for 2 processes, using point-to-point communication operations. You can assume that the matrix **A** is initially in process 0. Regarding the vector **w**, its initial content is not used and its correct final content can remain in any one of the processes. Justify the task assignment used.

**Solution:** We use the assignment:

$P_0 : T_1, T_2, T_5$

$P_1 : T_3, T_4, T_6$

Such an assignment maximizes parallelism, since the independent tasks are in different processes. In addition, communications are minimized by avoiding communicating the **A** matrix.

```
void func_par(double A[N][N], double w[N]) {
    double x[N], y[N], v[N], alpha;
    int rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    if (rank==0) {
        T1(A, x);
        MPI_Send(x, N, MPI_DOUBLE, 1, 0, MPI_COMM_WORLD);
        T2(A, x, y);
        MPI_Sendrecv(y, N, MPI_DOUBLE, 1, 0, v, N, MPI_DOUBLE, 1, 0,
                     MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        T5(A, y, v, &alpha);
        MPI_Send(&alpha, 1, MPI_DOUBLE, 1, 0, MPI_COMM_WORLD);
    } else if (rank==1) {
        MPI_Recv(x, N, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        T3(x, v);
        MPI_Sendrecv(v, N, MPI_DOUBLE, 0, 0, y, N, MPI_DOUBLE, 0, 0,
                     MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        T4(y, v, w);
        MPI_Recv(&alpha, 1, MPI_DOUBLE, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        T6(alpha, w);
    }
}
```

0.2 p.

- (c) Calculate the sequential cost and the parallel cost.

**Solution:** Sequential cost:

$$t(N) = 4 \cdot 4N^2 + 2 \cdot 3N^2 = 22N^2 \text{ flops}$$

Parallel cost:

$$t_a(N, 2) = 4N^2 + 4N^2 + 4N^2 + 4N^2 = 16N^2 \text{ flops}$$

$$t_c(N, 2) = 3(t_s + Nt_w) + (t_s + t_w) = 4t_s + (3N + 1)t_w \approx 4t_s + 3Nt_w$$

$$t(N, 2) = 16N^2 \text{ flops} + 4t_s + 3Nt_w$$