

## S2. Programming with OpenMP

J. M. Alonso, P. Alonso, F. Alvarruiz, I. Blanquer,  
J. Ibáñez, E. Ramos, J. E. Román

Departament de Sistemes Informàtics i Computació  
Universitat Politècnica de València

Year 2024/25



1

### Contents

- 1 Basic Concepts**
  - Programming Model
  - Simple Example
- 2 Loop Parallelization**
  - `parallel` for Directive
  - Variable Scope
  - Performance Improvement
- 3 Parallel Regions**
  - `parallel` Directive
  - Work Sharing
- 4 Synchronization**
  - Mutual Exclusion
  - Other Type of Synchronization

2

## Section 1

# Basic Concepts

- Programming Model
- Simple Example

3

## OpenMP Specification

De facto standard for shared memory programming

<http://www.openmp.org>

Specifications:

- Fortran: 1.0 (1997), 2.0 (2000)
- C/C++: 1.0 (1998), 2.0 (2002)
- Fortran/C/C++: 2.5 (2005), 3.0 (2008), 3.1 (2011), 4.0 (2013), 4.5 (2015), 5.0 (2018)

Previous experiences:

- ANSI X3H5 Standard (1994)
- HPF, CMFortran

4

## Programming Model

OpenMP programming is mainly based on **compiler directives**

### Example

```
void daxpy(int n, double a, double *x, double *y, double *z)
{
    int i;
    #pragma omp parallel for
    for (i=0; i<n; i++)
        z[i] = a*x[i] + y[i];
}
```

### Advantages

- It eases the adaptation (the compiler ignores #pragma)
- It enables incremental parallelization
- It enables compiler-based optimization

Additionally: functions (see `omp.h`) and environment variables

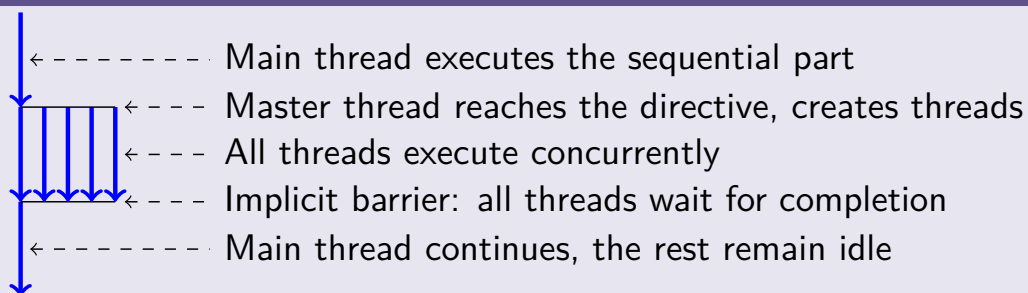
5

## Execution Model

OpenMP execution models follows the *fork-join* scheme

There are directives to create threads and sharing the work

### Scheme



Directives define **parallel regions**

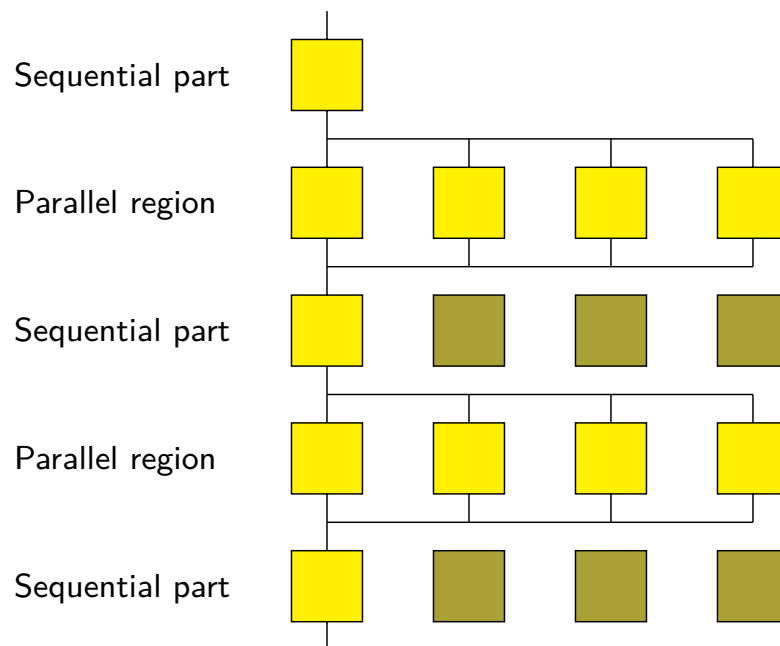
Other directives/clauses:

- Scope of variable: `private`, `shared`, `reduction`
- Synchronization: `critical`, `barrier`

6

## Execution Model - Threads

In OpenMP, idle threads are not destroyed, they remain ready for the next parallel region

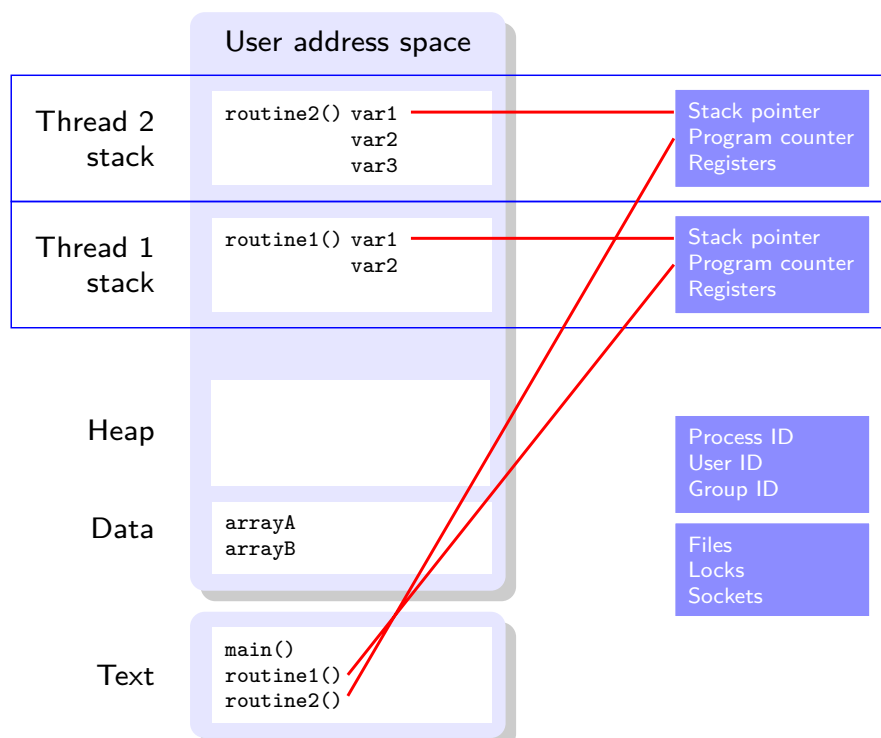


Threads created by a directive are called **team**

7

## Execution Model - Memory

Each thread has its own execution context (including the stack)



8

## Syntax

Directives:

```
#pragma omp <directive> [clause [...]]
```

Usage of functions:

```
#include <omp.h>
...
iam = omp_get_thread_num();
```

Conditional compilation: the `_OPENMP` macro contains the date of the supported OpenMP version, e.g. 201107

Compilation:

```
gcc> gcc -fopenmp prg-omp.c
sun> cc -xopenmp -x03 prg-omp.c
intel> icc -qopenmp prg-omp.c
```

9

## Simple Example

### Example

```
void daxpy(int n, double a, double *x, double *y, double *z)
{
    int i;
    #pragma omp parallel for
    for (i=0; i<n; i++)
        z[i] = a*x[i] + y[i];
}
```

- When `parallel` directive is reached, threads are created (if they have not been created previously)
- Loop iterations are shared among the threads
- By default, all the variables are shared, except for the loop variable (`i`) that is always private
- At the end, all threads synchronize

10

## Number of Threads and Thread Identifier

The number of threads can be specified:

- Using the `num_threads` clause
- Calling the `omp_set_num_threads()` function *before* the parallel region
- At run time, with `OMP_NUM_THREADS`

Useful functions:

- `omp_get_num_threads()`: returns the number of threads
- `omp_get_thread_num()`: returns the identifier of the thread (starting from 0, main thread is always 0)

```
omp_set_num_threads(3);
printf("threads before = %d\n",omp_get_num_threads());
#pragma omp parallel for
for (i=0; i<n; i++) {
    printf("threads = %d\n",omp_get_num_threads());
    printf("I am %d\n",omp_get_thread_num());
}
```

11

## Section 2

### Loop Parallelization

- `parallel for` Directive
- Variable Scope
- Performance Improvement

12

## The parallel for Directive

The loop below the directive is parallelized

C/C++

```
#pragma omp parallel for [clause [...]]
for (index=first; test_expr; increment_expr) {
    // loop body
}
```

OpenMP imposes restrictions to the loop type, for instance:

```
for (i=0; i<n && !found; i++)
    if (x[i]==elem) found=1;
```


13

## Example of parallel for

Consider a possible execution  
with 3 threads

Simple loop

```
#pragma omp parallel for
for (i=0; i<5; i++) {
    a[i] = calc(i);
}
```



Implicit barrier at the end of the parallel for construct

Variables:

- a: concurrent access, but not more than one thread accessing the same position
- i: different value in each thread → need a private copy

14

## Variable Scope

Variables are classified according to their scope:

- **Private**: each thread has a different replica
- **Shared**: all threads can read and write

Typical source of errors: choose an incorrect scope

---

The scope can be modified with clauses added to the directives:

- `private`, `shared`
- `reduction`
- `firstprivate`, `lastprivate`

15

## `private`, `shared`, `default`

If the scope of a variable is not specified, by default it is `shared`

Exceptions (`private`):

- Index variable of the parallelized loop
- Local variables of the called subroutines (except if they are declared `static`)
- Automatic variables declared inside the loop

Clause `default`

- `default(none)` forces to specify the scope of all variables

16



## private, shared

### private

```
sum = 0;
#pragma omp parallel for private(sum)
for (i=0; i<n; i++) {
    sum = sum + x[i]*x[i];
}
```

*Wrong:* after the loop, only the sum of the main thread is available  
- moreover, copies of each thread are not initialized

### shared

```
sum = 0;
#pragma omp parallel for shared(sum)
for (i=0; i<n; i++) {
    sum = sum + x[i]*x[i];
}
```

*Wrong:* race condition when reading/writing sum

17

## reduction

To perform reductions with commutative and associative operators  
(+, \*, -, &, |, ^, &&, ||, max, min)

### reduction(redn\_oper: var\_list)

```
sum = 0;
#pragma omp parallel for reduction(+:sum)
for (i=0; i<n; i++) {
    sum = sum + x[i]*x[i];
}
```

Each thread computes part of the sum, at the end all parts are combined in the total sum

It works as a private variable, but:

- At the end, the private values are combined
- Correctly initialized (to the neutral element of the operation)

18

## firstprivate, lastprivate

Private variables are created without an initial value and after the block its value is undefined

- **firstprivate**: initializes to the value of the main thread at the beginning of the block
- **lastprivate**: the value of the variable after the block is the one of the “last” iteration of the loop

### Example

```
alpha = 5.0;
#pragma omp parallel for firstprivate(alpha) lastprivate(i)
for (i=0; i<n; i++) {
    z[i] = alpha*x[i];
}
k = i;    /* i has value n */
```

Default behaviour tries to avoid unnecessary copies

19

## Guarantee Sufficient Work

Loop parallelization introduces an **overhead**: activation and deactivation of threads, synchronization

In simple loops, the overhead could be even larger than the compute time

### if clause

```
#pragma omp parallel for if(n>5000)
for (i=0; i<n; i++)
    z[i] = a*x[i] + y[i];
```

If the expression is false, the loop is executed sequentially

This clause can be also used to avoid data dependencies detected at execution time

20

## Nested Loops

We must put the directive before the loop to be parallelized

### Case 1

```
#pragma omp parallel for \
    private(j)
for (i=0; i<n; i++) {
    for (j=0; j<m; j++) {
        // loop body
    }
}
```

### Case 2

```
for (i=0; i<n; i++) {
    #pragma omp parallel for
    for (j=0; j<m; j++) {
        // loop body
    }
}
```

- In the first case, the iterations of *i* are shared; each thread executes the whole *j* loop
- In the second case, at each iteration of *i* threads are activated and deactivated; there are *n* synchronizations

21

## Nested Loops - Exchange

Usually recommended to parallelize the outer loop

- When data dependencies prevent from parallelizing the outer loop, **loop exchange** may be convenient

### Sequential code

```
for (j=1; j<n; j++)
    for (i=0; i<n; i++)
        a[i][j] = a[i][j] + a[i][j-1];
```

### Parallel code with exchanged loops

```
#pragma omp parallel for private(j)
for (i=0; i<n; i++)
    for (j=1; j<n; j++)
        a[i][j] = a[i][j] + a[i][j-1];
```

These changes may have an impact on cache use (locality)

22

## Scheduling

Ideally, all iterations cost the same and each thread gets assigned approximately the same number of iterations

In reality, a **workload imbalance** may appear, therefore reducing performance

---

In OpenMP it is possible to select the **scheduling**

Scheduling can be:

- Static: iterations are assigned to threads a priori
- Dynamic: assignment adapts to the current execution

The scheduling is done on the basis of contiguous ranges of iterations (*chunks*)

23

## Scheduling - schedule Clause

Syntax of the scheduling clause:

```
schedule(type[, chunk])
```

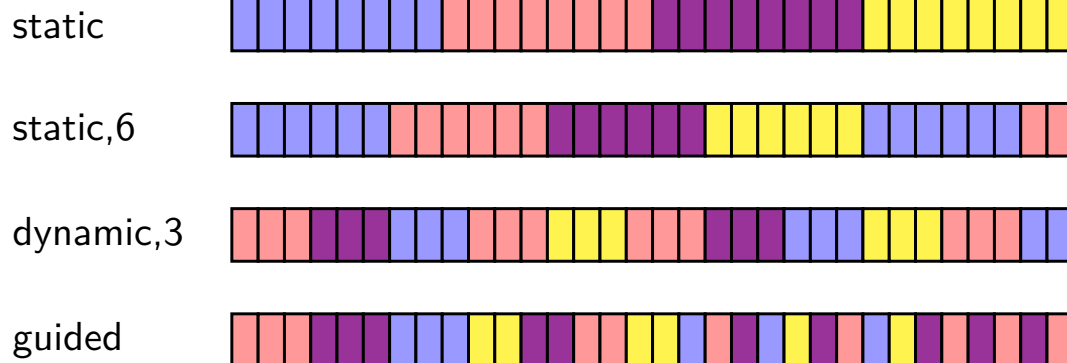
- static (without chunk): each thread receives an iteration range of similar size
- static (with chunk): cyclic assignment (*round-robin*) of ranges of size chunk
- dynamic (optional chunk, 1 by default): ranges are being assigned as required (*first-come, first-served*)
- guided (optional minimum chunk): same as dynamic but the size of the iteration range decreases exponentially ( $\propto n_{rem}/n_{threads}$ ) with the loop progress
- runtime: the scheduling is defined by the value of the environment variable OMP\_SCHEDULE

24

## Scheduling - Example

Example: loop of 32 iterations executed with 4 threads

```
$ OMP_NUM_THREADS=4 OMP_SCHEDULE=guided ./prog
```



25

### Section 3

## Parallel Regions

- parallel Directive
- Work Sharing

26

## The parallel Directive

The block below the directive is executed in a replicated way

C/C++

```
#pragma omp parallel [clause [clause ...]]
{
    // block
}
```

Some of the allowed clauses are: private, shared, default, reduction, if

Example - prints as many lines as threads

```
#pragma omp parallel private(myid)
{
    myid = omp_get_thread_num();
    printf("I am thread %d\n",myid);
}
```

27

## Work Sharing

Along with the replicated execution, it is often necessary to share the work among the threads

- Each thread works on a part of the data structure, or
- Each thread performs a different operation

---

Possible mechanisms for worksharing:

- Based on the thread identifier
- Parallel task queue
- Using OpenMP specific constructs

28

## Sharing Based on Thread Identifier

We use the following functions:

- `omp_get_num_threads()`: returns the number of threads
  - `omp_get_thread_num()`: returns the thread identifier
- to determine which part of the workload is done by each thread

### Example - thread identifiers

```
#pragma omp parallel private(myid)
{
    nthreads = omp_get_num_threads();
    myid = omp_get_thread_num();
    dowork(myid, nthreads);
}
```

29

## Sharing Using a Parallel Task Queue

A parallel task queue is a shared data structure containing a list of “tasks” to be performed

- Tasks can be processed concurrently
- Any task can be run by any thread

```
int get_next_task() {
    static int index = 0;
    int result;
    #pragma omp critical
    { if (index==MAXIDX) result=-1;
      else { index++; result=index; }
    }
    return result;
}
...
int myindex;
#pragma omp parallel private(myindex)
{ myindex = get_next_task();
  while (myindex>-1) {
      process_task(myindex);
      myindex = get_next_task();
  }
}
```

30

## Worksharing Constructs

The solutions mentioned before are quite primitive

- Programmer is in charge of splitting the workload
  - Obscure and complicated code in large programs
- 

OpenMP offers specific worksharing constructs

Three types:

- `for` construct to split iterations of loops
- Sections to distinguish different parts of the code
- Code to be executed by a single thread

Implicit barrier at the end of the block

31

## The `for` construct

Automatically distributes the iterations of a loop

### Example of shared loop

```
#pragma omp parallel
{
    ...
    #pragma omp for
    for (i=1; i<n; i++)
        b[i] = (a[i] + a[i-1]) / 2.0;
}
```

The loop iterations are not replicated but *shared* among the threads

`parallel` and `for` directives can be combined in one

32



## Loop Construct - nowait Clause

When several independent loops appear in the same parallel region, `nowait` removes the implicit barrier

### Loops without a barrier

```
void a8(int n, int m, float *a, float *b, float *y, float *z)
{
    int i;
    #pragma omp parallel
    {
        #pragma omp for nowait
        for (i=1; i<n; i++)
            b[i] = (a[i] + a[i-1]) / 2.0;

        #pragma omp for
        for (i=0; i<m; i++)
            y[i] = sqrt(z[i]);
    }
}
```

33

## The sections Construct

For independent pieces of code difficult to parallelize

- Individually they represent little work, or
- Each fragment is inherently sequential

It can also be combined with `parallel`

### Example of sections

```
#pragma omp parallel sections
{
    #pragma omp section
    Xaxis();
    #pragma omp section
    Yaxis();
    #pragma omp section
    Zaxis();
}
```

A thread may execute more than one section

Clauses: `private`, `first/lastprivate`, `reduction`, `nowait`

34

# The single Construct

Code fragments that must be executed by a single thread

## Example of single

```
#pragma omp parallel
{
    #pragma omp single
    printf("work1 starts\n");
    work1();

    #pragma omp single
    printf("work1 ends\n");

    #pragma omp single nowait
    printf("work1 ended, work2 starts\n");
    work2();
}
```

Some allowed clauses: private, firstprivate, nowait

35

## Section 4

# Synchronization

- Mutual Exclusion
- Other Type of Synchronization

36

## Race Condition (1)

The following example illustrates a race condition

### Find the maximum value

```
cur_max = -100000;  
#pragma omp parallel for  
for (i=0; i<n; i++) {  
    if (a[i] > cur_max) {  
        cur_max = a[i];  
    }  
}
```

Sequence with a wrong result:

Thread 0: reads a[i]=20, reads cur\_max=15

Thread 1: reads a[i]=16, reads cur\_max=15

Thread 0: checks a[i]>cur\_max, writes cur\_max=20

Thread 1: checks a[i]>cur\_max, writes cur\_max=16

37

## Race Condition (2)

There are cases where concurrent access does not produce a race condition

### Example of concurrent access without race condition

```
found = 0;  
#pragma omp parallel for  
for (i=0; i<n; i++) {  
    if (a[i] == value) {  
        found = 1;  
    }  
}
```

Even if several threads write at once, the result is correct

---

In general, synchronization mechanisms are needed:

- Mutual exclusion
- Other type of synchronization

38

## Mutual Exclusion

*Mutual exclusion* when accessing shared variables prevents any race condition

OpenMP provides three different constructs:

- Critical sections: `critical` directive
- Atomic operations: `atomic` directive
- Locks: `*_lock` routines

39

## The critical Directive (1)

In the previous example, access in mutual exclusion to variable `cur_max` prevents the race condition to happen

Find the maximum value, without race condition

```
cur_max = -100000;  
#pragma omp parallel for  
for (i=0; i<n; i++) {  
    #pragma omp critical  
    if (a[i] > cur_max) {  
        cur_max = a[i];  
    }  
}
```

When a thread reaches the `if` block (the critical section), it waits until no other thread is executing it at the same time

OpenMP guarantees **progress** (at least one waiting thread enters the critical section), but not **limited wait time**

40

## The critical Directive (2)

In practice, the previous example is executed sequentially

Considering that `cur_max` is never decreased, the following improvement can be introduced

### Improved maximum search

```
cur_max = -100000;
#pragma omp parallel for
for (i=0; i<n; i++) {
    if (a[i] > cur_max) {
        #pragma omp critical
        if (a[i] > cur_max)
            cur_max = a[i];
    }
}
```

The second `if` is required since `cur_max` is read outside the critical section

This solution enters the critical section less frequently

41

## Named critical Directive

By adding a name, we can have several unrelated critical sections

### Find the minimum and maximum values

```
cur_max = -100000;
cur_min = 100000;
#pragma omp parallel for
for (i=0; i<n; i++) {
    if (a[i] > cur_max) {
        #pragma omp critical (maxlock)
        if (a[i] > cur_max)
            cur_max = a[i];
    }
    if (a[i] < cur_min) {
        #pragma omp critical (minlock)
        if (a[i] < cur_min)
            cur_min = a[i];
    }
}
```

42

## The atomic Directive

Atomic load-modify-store operations

```
#pragma omp atomic
x <binop>= expr
```

```
#pragma omp atomic
x++, ++x, x--, --x
```

where <binop> can be +, \*, -, /, %, &, |, ^, <<, >>

### Example

```
#pragma omp parallel for shared(x, index, n)
for (i=0; i<n; i++) {
    #pragma omp atomic
    x[index[i]] += work1(i);
}
```

The code is much more efficient than using critical and enables updating the elements of x in parallel

43

## The barrier Directive

When reaching a barrier, threads wait for the rest to arrive

### Barrier example

```
#pragma omp parallel private(index)
{
    index = generate_next_index();
    while (index>0) {
        add_index(index);
        index = generate_next_index();
    }
    #pragma omp barrier
    index = get_next_index();
    while (index>0) {
        process_index(index);
        index = get_next_index();
    }
}
```

It is used to guarantee that a phase has been finished completely before proceeding to the next phase

44

## The ordered Directive

To make sure that a portion of the code of the iterations is executed in the original **sequential order**

### Example ordered

```
#pragma omp parallel for ordered
for (i=0; i<n; i++) {
    a[i] = ... /* complex computation */
    #pragma omp ordered
    fprintf(fd, "%d %g\n", i, a[i]);
}
```

Restrictions:

- If a parallel loop includes an ordered directive, the ordered clause should also be added to the loop
- Only one ordered section is allowed per iteration