

# Sesión 3

En esta tercera sesión aplicaremos regresión logística a algunas tareas de clasificación. Este cuaderno describe cómo hacerlo en iris con la implementación de regresión logística de sklearn. El objetivo principal de esta sesión es extender el ejemplo dado a otras tareas, tratando de minimizar el error de test.

# Regresión logística aplicada a iris

## Lectura del corpus y partición:

```
In [1]: import numpy as np; from sklearn.datasets import load_iris
        from sklearn.model_selection import train_test_split
        iris = load_iris(); X = iris.data.astype(np.float16); y = iris.target.astype(np.uint)
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=True, random_state=23)
```

## LogisticRegression: implementación de regresión logística en sklearn

```
In [2]: from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import accuracy_score
        clf = LogisticRegression(random_state=23).fit(X_train, y_train)
        y_test_pred = clf.predict(X_test)
        err_test = 1 - accuracy_score(y_test, clf.predict(X_test))
        print(f"Error de test: {err_test:.1%}")
```

Error de test: 0.0%

**Warnings:** sklearn es un poco "insistente" con los warnings; ignoraremos los avisos sobre convergencia

```
In [3]: import warnings; from sklearn.exceptions import ConvergenceWarning
        warnings.filterwarnings("ignore", category=ConvergenceWarning, module="sklearn")
```

**Solvers:** el parámetro `solver` de LogisticRegression permite elegir entre diferentes solvers (algoritmos de optimización)

```
In [4]: for solver in ['lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag', 'saga']:
        clf = LogisticRegression(random_state=23, solver=solver, max_iter=10000).fit(X_train, y_train)
        err_test = 1 - accuracy_score(y_test, clf.predict(X_test))
        print(f"Error de test después de entrenar con el solver {solver!s}: {err_test:.1%}")
```

```
Error de test después de entrenar con el solver lbfgs: 0.0%
Error de test después de entrenar con el solver liblinear: 3.3%
Error de test después de entrenar con el solver newton-cg: 0.0%
Error de test después de entrenar con el solver newton-cholesky: 3.3%
Error de test después de entrenar con el solver sag: 0.0%
Error de test después de entrenar con el solver saga: 0.0%
```

**Tolerancia:** el parámetro `tol` establece un umbral de tolerancia para acabar el entrenamiento (1e4 por defecto)

```
In [5]: for tol in (1e-4, 1e-2, 1, 1e2, 1e4):
        clf = LogisticRegression(tol=tol, random_state=23, max_iter=10000).fit(X_train, y_train)
        err_test = 1 - accuracy_score(y_test, clf.predict(X_test))
        print(f"Error de test con tolerancia {tol}: {err_test:.1%}")
```

```
Error de test con tolerancia 0.0001: 0.0%
Error de test con tolerancia 0.01: 3.3%
Error de test con tolerancia 1: 60.0%
Error de test con tolerancia 100.0: 60.0%
Error de test con tolerancia 10000.0: 60.0%
```

**Regularización:** el parámetro `C` (positivo, 1.0 por defecto) des-regulariza el criterio de entrenamiento

- **Posibilidad de subajuste:** con un valor próximo a cero (máxima regularización)
- **Posibilidad de sobreajuste:** con un valor positivo muy alto (mínima regularización)

```
In [6]: for C in (1e-2, 1e-1, 1, 1e1, 1e2):  
        clf = LogisticRegression(C=C, random_state=23, max_iter=10000).fit(X_train, y_train)  
        err_test = 1 - accuracy_score(y_test, clf.predict(X_test))  
        print(f"Error de test con C {C:g}: {err_test:.1%}")
```

```
Error de test con C 0.01: 6.7%  
Error de test con C 0.1: 3.3%  
Error de test con C 1: 0.0%  
Error de test con C 10: 3.3%  
Error de test con C 100: 3.3%
```

**Early stopping:** ahorramos cálculo y evitamos sobre-entrenamiento ("regularizamos") acabando pronto (en pocas iteraciones)

```
In [7]: for max_iter in (10, 20, 50, 100):  
        clf = LogisticRegression(random_state=23, max_iter=max_iter).fit(X_train, y_train)  
        err_test = 1 - accuracy_score(y_test, clf.predict(X_test))  
        print(f"Error de test con max_iter {max_iter}: {err_test:.1%}")
```

```
Error de test con max_iter 10: 0.0%  
Error de test con max_iter 20: 3.3%  
Error de test con max_iter 50: 0.0%  
Error de test con max_iter 100: 0.0%
```

# Regresión logística aplicada a openml

```
In [1]: import warnings; warnings.filterwarnings("ignore"); import numpy as np
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
In [2]: def err_eval(data_id):
    X, y = fetch_openml(data_id=data_id, return_X_y=True, as_frame=False, parser="liac-arff")
    mask = ~np.isnan(X).any(axis=1); X = X[mask, :]; y = y[mask]
    if X.shape[0] < 10: return(1.0)
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=True, random_state=23)
    clf = LogisticRegression(random_state=23).fit(X_train, y_train)
    return(1 - accuracy_score(y_test, clf.predict(X_test)))
```

```
In [3]: import openml
for sid in (99, 334):
    benchmark_suite = openml.study.get_suite(suite_id=sid)
    df = openml.datasets.list_datasets(data_id=benchmark_suite.data, output_format='dataframe')
    for did, name, C in zip(df['did'], df['name'], df['NumberOfClasses']):
        err = err_eval(did)
        print(f"sid: {sid:5d} did: {did:5d} C: {C:5.0f} err: {err:7.1%} name: {name:s}")
```

sid:	99	did:	3	C:	2	err:	5.5%	name:	kr-vs-kp
sid:	99	did:	6	C:	26	err:	23.7%	name:	letter
sid:	99	did:	11	C:	3	err:	8.0%	name:	balance-scale
sid:	99	did:	12	C:	10	err:	4.2%	name:	mfeat-factors
sid:	99	did:	14	C:	10	err:	21.0%	name:	mfeat-fourier
sid:	99	did:	15	C:	2	err:	2.2%	name:	breast-w
sid:	99	did:	16	C:	10	err:	4.5%	name:	mfeat-karhunen
sid:	99	did:	18	C:	10	err:	46.8%	name:	mfeat-morphological
sid:	99	did:	22	C:	10	err:	18.2%	name:	mfeat-zernike
sid:	99	did:	23	C:	3	err:	51.2%	name:	cmc
sid:	99	did:	28	C:	10	err:	3.3%	name:	optdigits
sid:	99	did:	29	C:	2	err:	13.7%	name:	credit-approval
sid:	99	did:	31	C:	2	err:	25.5%	name:	credit-g
sid:	99	did:	32	C:	10	err:	6.0%	name:	pendigits
sid:	99	did:	37	C:	2	err:	17.5%	name:	diabetes
sid:	99	did:	38	C:	2	err:	100.0%	name:	sick
sid:	99	did:	44	C:	2	err:	7.3%	name:	spambase
sid:	99	did:	46	C:	3	err:	6.7%	name:	splice
sid:	99	did:	50	C:	2	err:	29.2%	name:	tic-tac-toe
sid:	99	did:	54	C:	4	err:	29.4%	name:	vehicle
sid:	99	did:	151	C:	2	err:	24.9%	name:	electricity
sid:	99	did:	182	C:	6	err:	14.2%	name:	satimage
sid:	99	did:	188	C:	5	err:	60.5%	name:	eucalyptus
sid:	99	did:	300	C:	26	err:	4.5%	name:	isolet
sid:	99	did:	307	C:	11	err:	39.4%	name:	vowel
sid:	99	did:	458	C:	4	err:	1.2%	name:	analcata_data_authorship
sid:	99	did:	469	C:	6	err:	81.2%	name:	analcata_data_dmft
sid:	99	did:	554	C:	10	err:	8.0%	name:	mnist_784
sid:	99	did:	1049	C:	2	err:	12.0%	name:	pc4
sid:	99	did:	1050	C:	2	err:	8.6%	name:	pc3
sid:	99	did:	1053	C:	2	err:	19.5%	name:	jml
sid:	99	did:	1063	C:	2	err:	13.3%	name:	kc2
sid:	99	did:	1067	C:	2	err:	17.5%	name:	kc1
sid:	99	did:	1068	C:	2	err:	5.9%	name:	pc1
sid:	99	did:	1461	C:	2	err:	11.5%	name:	bank-marketing
sid:	99	did:	1462	C:	2	err:	1.5%	name:	banknote-authentication
sid:	99	did:	1464	C:	2	err:	27.3%	name:	blood-transfusion-service-center
sid:	99	did:	1468	C:	9	err:	6.0%	name:	cnae-9
sid:	99	did:	1475	C:	6	err:	53.7%	name:	first-order-theorem-proving
sid:	99	did:	1478	C:	6	err:	1.8%	name:	har
sid:	99	did:	1480	C:	2	err:	29.9%	name:	ilpd
sid:	99	did:	1485	C:	2	err:	42.9%	name:	madelon
sid:	99	did:	1486	C:	2	err:	5.2%	name:	nomao
sid:	99	did:	1487	C:	2	err:	5.9%	name:	ozone-level-8hr
sid:	99	did:	1489	C:	2	err:	25.3%	name:	phoneme

sid:	99	did:	1494	C:	2	err:	16.6%	name:	qsar-biodeg
sid:	99	did:	1497	C:	4	err:	29.6%	name:	wall-robot-navigation
sid:	99	did:	1501	C:	10	err:	9.4%	name:	semeion
sid:	99	did:	1510	C:	2	err:	3.5%	name:	wdbc
sid:	99	did:	1590	C:	2	err:	20.8%	name:	adult
sid:	99	did:	4134	C:	2	err:	24.0%	name:	Bioresponse
sid:	99	did:	4534	C:	2	err:	10.2%	name:	PhishingWebsites
sid:	99	did:	4538	C:	5	err:	54.6%	name:	GesturePhaseSegmentationProcessed
sid:	99	did:	6332	C:	2	err:	28.6%	name:	cylinder-bands
sid:	99	did:	23381	C:	2	err:	50.0%	name:	dresses-sales
sid:	99	did:	23517	C:	2	err:	47.4%	name:	numera128.6
sid:	99	did:	40499	C:	11	err:	2.1%	name:	texture
sid:	99	did:	40668	C:	3	err:	33.4%	name:	connect-4
sid:	99	did:	40670	C:	3	err:	5.2%	name:	dna
sid:	99	did:	40701	C:	2	err:	13.4%	name:	churn
sid:	99	did:	40923	C:	46	err:	27.7%	name:	Devnagari-Script
sid:	99	did:	40927	C:	10	err:	58.9%	name:	CIFAR_10
sid:	99	did:	40966	C:	8	err:	6.3%	name:	MiceProtein
sid:	99	did:	40975	C:	4	err:	15.6%	name:	car
sid:	99	did:	40978	C:	2	err:	3.2%	name:	Internet-Advertisements
sid:	99	did:	40979	C:	10	err:	3.2%	name:	mfeat-pixel
sid:	99	did:	40982	C:	7	err:	58.4%	name:	steel-plates-fault
sid:	99	did:	40983	C:	2	err:	3.5%	name:	wilt
sid:	99	did:	40984	C:	7	err:	7.8%	name:	segment
sid:	99	did:	40994	C:	2	err:	5.6%	name:	climate-model-simulation-crashes
sid:	99	did:	40996	C:	10	err:	14.7%	name:	Fashion-MNIST
sid:	99	did:	41027	C:	3	err:	31.8%	name:	jungle_chess_2pcs_raw_endgame_complete
sid:	334	did:	44156	C:	2	err:	26.6%	name:	electricity
sid:	334	did:	44157	C:	2	err:	46.8%	name:	eye_movements
sid:	334	did:	44159	C:	2	err:	24.5%	name:	covertime
sid:	334	did:	45035	C:	2	err:	36.8%	name:	albert
sid:	334	did:	45036	C:	2	err:	37.7%	name:	default-of-credit-card-clients
sid:	334	did:	45038	C:	2	err:	41.6%	name:	road-safety
sid:	334	did:	45039	C:	2	err:	30.2%	name:	compas-two-years

```
In [4]: import openml
for sid in (271, ):
    benchmark_suite = openml.study.get_suite(suite_id=sid)
    df = openml.datasets.list_datasets(data_id=benchmark_suite.data, output_format='dataframe')
    for did, name, C in zip(df['did'], df['name'], df['NumberOfClasses']):
        if did == 41147: continue;
        err = err_eval(did)
        print(f"sid: {sid:5d}  did: {did:5d}  C: {C:5.0f}  err: {err:7.1%}  name: {name:s}")
```



sid:	271	did:	3	C:	2	err:	5.5%	name:	kr-vs-kp
sid:	271	did:	12	C:	10	err:	4.2%	name:	mfeat-factors
sid:	271	did:	23	C:	3	err:	51.2%	name:	cmc
sid:	271	did:	31	C:	2	err:	25.5%	name:	credit-g
sid:	271	did:	54	C:	4	err:	29.4%	name:	vehicle
sid:	271	did:	181	C:	10	err:	45.1%	name:	yeast
sid:	271	did:	188	C:	5	err:	60.5%	name:	eucalyptus
sid:	271	did:	1049	C:	2	err:	12.0%	name:	pc4
sid:	271	did:	1067	C:	2	err:	17.5%	name:	kc1
sid:	271	did:	1111	C:	2	err:	100.0%	name:	KDDCup09_appetency
sid:	271	did:	1169	C:	2	err:	42.1%	name:	airlines
sid:	271	did:	1457	C:	50	err:	37.3%	name:	amazon-commerce-reviews
sid:	271	did:	1461	C:	2	err:	11.5%	name:	bank-marketing
sid:	271	did:	1464	C:	2	err:	27.3%	name:	blood-transfusion-service-center
sid:	271	did:	1468	C:	9	err:	6.0%	name:	cnae-9
sid:	271	did:	1475	C:	6	err:	53.7%	name:	first-order-theorem-proving
sid:	271	did:	1486	C:	2	err:	5.2%	name:	nomao
sid:	271	did:	1487	C:	2	err:	5.9%	name:	ozone-level-8hr
sid:	271	did:	1489	C:	2	err:	25.3%	name:	phoneme
sid:	271	did:	1494	C:	2	err:	16.6%	name:	qsar-biodeg
sid:	271	did:	1515	C:	20	err:	13.0%	name:	micro-mass
sid:	271	did:	1590	C:	2	err:	20.8%	name:	adult
sid:	271	did:	1596	C:	7	err:	38.3%	name:	coverttype
sid:	271	did:	4134	C:	2	err:	24.0%	name:	Bioresponse
sid:	271	did:	4135	C:	2	err:	5.6%	name:	Amazon_employee_access
sid:	271	did:	4534	C:	2	err:	10.2%	name:	PhishingWebsites
sid:	271	did:	4538	C:	5	err:	54.6%	name:	GesturePhaseSegmentationProcessed
sid:	271	did:	4541	C:	3	err:	46.0%	name:	Diabetes130US
sid:	271	did:	23517	C:	2	err:	47.4%	name:	numera128.6
sid:	271	did:	40498	C:	7	err:	54.4%	name:	wine-quality-white
sid:	271	did:	40668	C:	3	err:	33.4%	name:	connect-4
sid:	271	did:	40670	C:	3	err:	5.2%	name:	dna
sid:	271	did:	40685	C:	7	err:	7.8%	name:	shuttle
sid:	271	did:	40701	C:	2	err:	13.4%	name:	churn
sid:	271	did:	40900	C:	2	err:	0.8%	name:	Satellite
sid:	271	did:	40975	C:	4	err:	15.6%	name:	car
sid:	271	did:	40978	C:	2	err:	3.2%	name:	Internet-Advertisements
sid:	271	did:	40981	C:	2	err:	16.7%	name:	Australian
sid:	271	did:	40982	C:	7	err:	58.4%	name:	steel-plates-fault
sid:	271	did:	40983	C:	2	err:	3.5%	name:	wilt
sid:	271	did:	40984	C:	7	err:	7.8%	name:	segment
sid:	271	did:	40996	C:	10	err:	14.7%	name:	Fashion-MNIST
sid:	271	did:	41027	C:	3	err:	31.8%	name:	jungle_chess_2pcs_raw_endgame_complete
sid:	271	did:	41138	C:	2	err:	2.6%	name:	APSFailure
sid:	271	did:	41142	C:	2	err:	32.8%	name:	christine

sid:	271	did:	41143	C:	2	err:	19.8%	name:	jasmine
sid:	271	did:	41144	C:	2	err:	46.3%	name:	madeline
sid:	271	did:	41145	C:	2	err:	31.2%	name:	philippine
sid:	271	did:	41146	C:	2	err:	14.2%	name:	sylvine
sid:	271	did:	41150	C:	2	err:	14.8%	name:	MiniBooNE
sid:	271	did:	41156	C:	2	err:	17.1%	name:	ada
sid:	271	did:	41157	C:	2	err:	5.0%	name:	arcene
sid:	271	did:	41158	C:	2	err:	17.4%	name:	gina
sid:	271	did:	41159	C:	2	err:	33.3%	name:	guillermo
sid:	271	did:	41161	C:	2	err:	1.1%	name:	riccardo
sid:	271	did:	41162	C:	2	err:	5.0%	name:	kick
sid:	271	did:	41163	C:	5	err:	8.1%	name:	dilbert
sid:	271	did:	41164	C:	7	err:	31.6%	name:	fabert
sid:	271	did:	41165	C:	10	err:	63.2%	name:	robert
sid:	271	did:	41166	C:	10	err:	46.9%	name:	volkert
sid:	271	did:	41167	C:	355	err:	98.9%	name:	dionis
sid:	271	did:	41168	C:	4	err:	39.6%	name:	jannis
sid:	271	did:	41169	C:	100	err:	76.3%	name:	helena
sid:	271	did:	42732	C:	2	err:	12.2%	name:	sf-police-incidents
sid:	271	did:	42733	C:	2	err:	16.5%	name:	Click_prediction_small
sid:	271	did:	42734	C:	3	err:	26.5%	name:	okcupid-stem
sid:	271	did:	42742	C:	2	err:	4.6%	name:	porto-seguro
sid:	271	did:	42746	C:	23	err:	3.2%	name:	KDDCup99
sid:	271	did:	42769	C:	2	err:	35.9%	name:	Higgs
sid:	271	did:	43072	C:	2	err:	100.0%	name:	KDDCup09-Upselling