

# Session 3

In this third session we will apply logistic regression to some classification tasks. This notebook describes how to do it in Iris with the sklearn implementation of logistic regression. The main objective of this session is to extend the example given to other tasks, trying to minimize test error.

## Logistic regression applied to Iris

**Reading and partitioning the dataset:**

```
In [1]: import numpy as np; from sklearn.datasets import load_iris
        from sklearn.model_selection import train_test_split
        iris = load_iris(); X = iris.data.astype(np.float16); y = iris.target.astype(np.uint)
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=True, random_state=23)
```

**LogisticRegression:** implementation of logistic regression in sklearn

```
In [2]: from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import accuracy_score
        clf = LogisticRegression(random_state=23).fit(X_train, y_train)
        y_test_pred = clf.predict(X_test)
        err_test = 1 - accuracy_score(y_test, clf.predict(X_test))
        print(f"Test error: {err_test:.1%}")
```

Test error: 0.0%

**Warnings:** sklearn is a bit "picky" with warnings; we will ignore the warnings about convergence

```
In [3]: import warnings; from sklearn.exceptions import ConvergenceWarning
        warnings.filterwarnings("ignore", category=ConvergenceWarning, module="sklearn")
```

**Solvers:** the `solver` parameter of LogisticRegression allows you to choose between different solvers (optimization algorithms)

```
In [4]: for solver in ['lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky', 'sag', 'saga']:
        clf = LogisticRegression(random_state=23, solver=solver, max_iter=10000).fit(X_train, y_train)
        err_test = 1 - accuracy_score(y_test, clf.predict(X_test))
        print(f"Test error after training with the solver {solver!s}: {err_test:.1%}")
```

```
Test error after training with the solver lbfgs: 0.0%
Test error after training with the solver liblinear: 3.3%
Test error after training with the solver newton-cg: 0.0%
Test error after training with the solver newton-cholesky: 3.3%
Test error after training with the solver sag: 0.0%
Test error after training with the solver saga: 0.0%
```

**Tolerance:** the `tol` parameter sets a tolerance threshold to end training (1e4 by default)

```
In [5]: for tol in (1e-4, 1e-2, 1, 1e2, 1e4):
        clf = LogisticRegression(tol=tol, random_state=23, max_iter=10000).fit(X_train, y_train)
        err_test = 1 - accuracy_score(y_test, clf.predict(X_test))
        print(f"Test error with tolerance {tol}: {err_test:.1%}")
```

```
Test error with tolerance 0.0001: 0.0%
Test error with tolerance 0.01: 3.3%
Test error with tolerance 1: 60.0%
Test error with tolerance 100.0: 60.0%
Test error with tolerance 10000.0: 60.0%
```

**Regularization:** parameter `C` (positive, 1.0 by default) de-regularizes the training criterion

- **Possibility of under-adjustment:** with a value close to zero (maximum regularization)
- **Possibility of overfitting:** with a very high positive value (minimum regularization)

```
In [6]: for C in (1e-2, 1e-1, 1, 1e1, 1e2):  
        clf = LogisticRegression(C=C, random_state=23, max_iter=10000).fit(X_train, y_train)  
        err_test = 1 - accuracy_score(y_test, clf.predict(X_test))  
        print(f"Test error with C {C:g}: {err_test:.1%}")
```

```
Test error with C 0.01: 6.7%  
Test error with C 0.1: 3.3%  
Test error with C 1: 0.0%  
Test error with C 10: 3.3%  
Test error with C 100: 3.3%
```

**Early stopping:** saving computation and avoiding over-training ("regularize") by finishing earlier (in a few iterations)

```
In [7]: for max_iter in (10, 20, 50, 100):  
        clf = LogisticRegression(random_state=23, max_iter=max_iter).fit(X_train, y_train)  
        err_test = 1 - accuracy_score(y_test, clf.predict(X_test))  
        print(f"Test error with max_iter {max_iter}: {err_test:.1%}")
```

```
Test error with max_iter 10: 0.0%  
Test error with max_iter 20: 3.3%  
Test error with max_iter 50: 0.0%  
Test error with max_iter 100: 0.0%
```

# Logistic regression applied to openml

```
In [1]: import warnings; warnings.filterwarnings("ignore"); import numpy as np
from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

```
In [2]: def err_eval(data_id):
    X, y = fetch_openml(data_id=data_id, return_X_y=True, as_frame=False, parser="liac-arff")
    mask = ~np.isnan(X).any(axis=1); X = X[mask, :]; y = y[mask]
    if X.shape[0] < 10: return(1.0)
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=True, random_state=23)
    clf = LogisticRegression(random_state=23).fit(X_train, y_train)
    return(1 - accuracy_score(y_test, clf.predict(X_test)))
```

```
In [3]: import openml
        for sid in (99, 334):
            benchmark_suite = openml.study.get_suite(suite_id=sid)
            df = openml.datasets.list_datasets(data_id=benchmark_suite.data, output_format='dataframe')
            for did, name, C in zip(df['did'], df['name'], df['NumberOfClasses']):
                err = err_eval(did)
                print(f"sid: {sid:5d}  did: {did:5d}  C: {C:5.0f}  err: {err:7.1%}  name: {name:s}")
```

sid:	99	did:	3	C:	2	err:	5.5%	name:	kr-vs-kp
sid:	99	did:	6	C:	26	err:	23.7%	name:	letter
sid:	99	did:	11	C:	3	err:	8.0%	name:	balance-scale
sid:	99	did:	12	C:	10	err:	4.2%	name:	mfeat-factors
sid:	99	did:	14	C:	10	err:	21.0%	name:	mfeat-fourier
sid:	99	did:	15	C:	2	err:	2.2%	name:	breast-w
sid:	99	did:	16	C:	10	err:	4.5%	name:	mfeat-karhunen
sid:	99	did:	18	C:	10	err:	46.8%	name:	mfeat-morphological
sid:	99	did:	22	C:	10	err:	18.2%	name:	mfeat-zernike
sid:	99	did:	23	C:	3	err:	50.8%	name:	cmc
sid:	99	did:	28	C:	10	err:	3.3%	name:	optdigits
sid:	99	did:	29	C:	2	err:	13.7%	name:	credit-approval
sid:	99	did:	31	C:	2	err:	25.5%	name:	credit-g
sid:	99	did:	32	C:	10	err:	6.0%	name:	pendigits
sid:	99	did:	37	C:	2	err:	16.9%	name:	diabetes
sid:	99	did:	38	C:	2	err:	100.0%	name:	sick
sid:	99	did:	44	C:	2	err:	7.4%	name:	spambase
sid:	99	did:	46	C:	3	err:	6.7%	name:	splice
sid:	99	did:	50	C:	2	err:	29.2%	name:	tic-tac-toe
sid:	99	did:	54	C:	4	err:	29.4%	name:	vehicle
sid:	99	did:	151	C:	2	err:	24.9%	name:	electricity
sid:	99	did:	182	C:	6	err:	14.2%	name:	satimage
sid:	99	did:	188	C:	5	err:	60.5%	name:	eucalyptus
sid:	99	did:	300	C:	26	err:	4.5%	name:	isolet
sid:	99	did:	307	C:	11	err:	39.4%	name:	vowel
sid:	99	did:	458	C:	4	err:	1.2%	name:	analcata_data_authorship
sid:	99	did:	469	C:	6	err:	81.2%	name:	analcata_data_dmft
sid:	99	did:	554	C:	10	err:	8.0%	name:	mnist_784
sid:	99	did:	1049	C:	2	err:	11.6%	name:	pc4
sid:	99	did:	1050	C:	2	err:	8.9%	name:	pc3
sid:	99	did:	1053	C:	2	err:	19.5%	name:	jml
sid:	99	did:	1063	C:	2	err:	13.3%	name:	kc2
sid:	99	did:	1067	C:	2	err:	16.8%	name:	kc1
sid:	99	did:	1068	C:	2	err:	5.9%	name:	pc1
sid:	99	did:	1461	C:	2	err:	11.5%	name:	bank-marketing
sid:	99	did:	1462	C:	2	err:	1.5%	name:	banknote-authentication
sid:	99	did:	1464	C:	2	err:	27.3%	name:	blood-transfusion-service-center
sid:	99	did:	1468	C:	9	err:	6.0%	name:	cnae-9
sid:	99	did:	1475	C:	6	err:	53.7%	name:	first-order-theorem-proving
sid:	99	did:	1478	C:	6	err:	1.8%	name:	har
sid:	99	did:	1480	C:	2	err:	29.9%	name:	ilpd
sid:	99	did:	1485	C:	2	err:	43.7%	name:	madelon
sid:	99	did:	1486	C:	2	err:	5.2%	name:	nomao
sid:	99	did:	1487	C:	2	err:	5.3%	name:	ozone-level-8hr
sid:	99	did:	1489	C:	2	err:	25.3%	name:	phoneme

sid:	99	did:	1494	C:	2	err:	16.6%	name:	qsar-biodeg
sid:	99	did:	1497	C:	4	err:	29.6%	name:	wall-robot-navigation
sid:	99	did:	1501	C:	10	err:	9.4%	name:	semeion
sid:	99	did:	1510	C:	2	err:	3.5%	name:	wdbc
sid:	99	did:	1590	C:	2	err:	20.8%	name:	adult
sid:	99	did:	4134	C:	2	err:	24.0%	name:	Bioresponse
sid:	99	did:	4534	C:	2	err:	10.2%	name:	PhishingWebsites
sid:	99	did:	4538	C:	5	err:	54.6%	name:	GesturePhaseSegmentationProcessed
sid:	99	did:	6332	C:	2	err:	28.6%	name:	cylinder-bands
sid:	99	did:	23381	C:	2	err:	50.0%	name:	dresses-sales
sid:	99	did:	23517	C:	2	err:	47.4%	name:	numera128.6
sid:	99	did:	40499	C:	11	err:	2.1%	name:	texture
sid:	99	did:	40668	C:	3	err:	33.4%	name:	connect-4
sid:	99	did:	40670	C:	3	err:	5.2%	name:	dna
sid:	99	did:	40701	C:	2	err:	13.4%	name:	churn
sid:	99	did:	40923	C:	46	err:	27.7%	name:	Devnagari-Script
sid:	99	did:	40927	C:	10	err:	58.9%	name:	CIFAR_10
sid:	99	did:	40966	C:	8	err:	6.3%	name:	MiceProtein
sid:	99	did:	40975	C:	4	err:	15.6%	name:	car
sid:	99	did:	40978	C:	2	err:	3.2%	name:	Internet-Advertisements
sid:	99	did:	40979	C:	10	err:	3.2%	name:	mfeat-pixel
sid:	99	did:	40982	C:	7	err:	58.4%	name:	steel-plates-fault
sid:	99	did:	40983	C:	2	err:	3.5%	name:	wilt
sid:	99	did:	40984	C:	7	err:	7.8%	name:	segment
sid:	99	did:	40994	C:	2	err:	5.6%	name:	climate-model-simulation-crashes
sid:	99	did:	40996	C:	10	err:	14.7%	name:	Fashion-MNIST
sid:	99	did:	41027	C:	3	err:	31.8%	name:	jungle_chess_2pcs_raw_endgame_complete
sid:	334	did:	44156	C:	2	err:	26.6%	name:	electricity
sid:	334	did:	44157	C:	2	err:	47.9%	name:	eye_movements
sid:	334	did:	44159	C:	2	err:	24.5%	name:	coverttype
sid:	334	did:	45035	C:	2	err:	36.7%	name:	albert
sid:	334	did:	45036	C:	2	err:	37.7%	name:	default-of-credit-card-clients
sid:	334	did:	45038	C:	2	err:	42.0%	name:	road-safety
sid:	334	did:	45039	C:	2	err:	30.2%	name:	compas-two-years

```
In [4]: import openml
for sid in (271, ):
    benchmark_suite = openml.study.get_suite(suite_id=sid)
    df = openml.datasets.list_datasets(data_id=benchmark_suite.data, output_format='dataframe')
    for did, name, C in zip(df['did'], df['name'], df['NumberOfClasses']):
        if did == 41147: continue;
        err = err_eval(did)
        print(f"sid: {sid:5d}  did: {did:5d}  C: {C:5.0f}  err: {err:7.1%}  name: {name:s}")
```



sid:	271	did:	3	C:	2	err:	5.5%	name:	kr-vs-kp
sid:	271	did:	12	C:	10	err:	4.2%	name:	mfeat-factors
sid:	271	did:	23	C:	3	err:	50.8%	name:	cmc
sid:	271	did:	31	C:	2	err:	25.5%	name:	credit-g
sid:	271	did:	54	C:	4	err:	29.4%	name:	vehicle
sid:	271	did:	181	C:	10	err:	45.1%	name:	yeast
sid:	271	did:	188	C:	5	err:	60.5%	name:	eucalyptus
sid:	271	did:	1049	C:	2	err:	11.6%	name:	pc4
sid:	271	did:	1067	C:	2	err:	16.8%	name:	kc1
sid:	271	did:	1111	C:	2	err:	100.0%	name:	KDDCup09_appetency
sid:	271	did:	1169	C:	2	err:	42.1%	name:	airlines
sid:	271	did:	1457	C:	50	err:	37.3%	name:	amazon-commerce-reviews
sid:	271	did:	1461	C:	2	err:	11.5%	name:	bank-marketing
sid:	271	did:	1464	C:	2	err:	27.3%	name:	blood-transfusion-service-center
sid:	271	did:	1468	C:	9	err:	6.0%	name:	cnae-9
sid:	271	did:	1475	C:	6	err:	53.7%	name:	first-order-theorem-proving
sid:	271	did:	1486	C:	2	err:	5.2%	name:	nomao
sid:	271	did:	1487	C:	2	err:	5.3%	name:	ozone-level-8hr
sid:	271	did:	1489	C:	2	err:	25.3%	name:	phoneme
sid:	271	did:	1494	C:	2	err:	16.6%	name:	qsar-biodeg
sid:	271	did:	1515	C:	20	err:	13.9%	name:	micro-mass
sid:	271	did:	1590	C:	2	err:	20.8%	name:	adult
sid:	271	did:	1596	C:	7	err:	38.3%	name:	coverttype
sid:	271	did:	4134	C:	2	err:	24.0%	name:	Bioresponse
sid:	271	did:	4135	C:	2	err:	5.6%	name:	Amazon_employee_access
sid:	271	did:	4534	C:	2	err:	10.2%	name:	PhishingWebsites
sid:	271	did:	4538	C:	5	err:	54.6%	name:	GesturePhaseSegmentationProcessed
sid:	271	did:	4541	C:	3	err:	46.0%	name:	Diabetes130US
sid:	271	did:	23517	C:	2	err:	47.4%	name:	numera128.6
sid:	271	did:	40498	C:	7	err:	54.4%	name:	wine-quality-white
sid:	271	did:	40668	C:	3	err:	33.4%	name:	connect-4
sid:	271	did:	40670	C:	3	err:	5.2%	name:	dna
sid:	271	did:	40685	C:	7	err:	7.7%	name:	shuttle
sid:	271	did:	40701	C:	2	err:	13.4%	name:	churn
sid:	271	did:	40900	C:	2	err:	0.8%	name:	Satellite
sid:	271	did:	40975	C:	4	err:	15.6%	name:	car
sid:	271	did:	40978	C:	2	err:	3.2%	name:	Internet-Advertisements
sid:	271	did:	40981	C:	2	err:	16.7%	name:	Australian
sid:	271	did:	40982	C:	7	err:	58.4%	name:	steel-plates-fault
sid:	271	did:	40983	C:	2	err:	3.5%	name:	wilt
sid:	271	did:	40984	C:	7	err:	7.8%	name:	segment
sid:	271	did:	40996	C:	10	err:	14.7%	name:	Fashion-MNIST
sid:	271	did:	41027	C:	3	err:	31.8%	name:	jungle_chess_2pcs_raw_endgame_complete
sid:	271	did:	41138	C:	2	err:	2.6%	name:	APSFailure
sid:	271	did:	41142	C:	2	err:	33.2%	name:	christine

sid:	271	did:	41143	C:	2	err:	19.8%	name:	jasmine
sid:	271	did:	41144	C:	2	err:	45.5%	name:	madeline
sid:	271	did:	41145	C:	2	err:	31.2%	name:	philippine
sid:	271	did:	41146	C:	2	err:	14.0%	name:	sylvine
sid:	271	did:	41150	C:	2	err:	14.3%	name:	MiniBooNE
sid:	271	did:	41156	C:	2	err:	17.1%	name:	ada
sid:	271	did:	41157	C:	2	err:	5.0%	name:	arcene
sid:	271	did:	41158	C:	2	err:	17.4%	name:	gina
sid:	271	did:	41159	C:	2	err:	33.3%	name:	guillermo
sid:	271	did:	41161	C:	2	err:	1.1%	name:	riccardo
sid:	271	did:	41162	C:	2	err:	5.0%	name:	kick
sid:	271	did:	41163	C:	5	err:	8.3%	name:	dilbert
sid:	271	did:	41164	C:	7	err:	31.6%	name:	fabert
sid:	271	did:	41165	C:	10	err:	63.2%	name:	robert
sid:	271	did:	41166	C:	10	err:	46.9%	name:	volkert
sid:	271	did:	41167	C:	355	err:	98.7%	name:	dionis
sid:	271	did:	41168	C:	4	err:	39.6%	name:	jannis
sid:	271	did:	41169	C:	100	err:	76.3%	name:	helena
sid:	271	did:	42732	C:	2	err:	12.2%	name:	sf-police-incidents
sid:	271	did:	42733	C:	2	err:	16.5%	name:	Click_prediction_small
sid:	271	did:	42734	C:	3	err:	26.5%	name:	okcupid-stem
sid:	271	did:	42742	C:	2	err:	4.6%	name:	porto-seguro
sid:	271	did:	42746	C:	23	err:	2.9%	name:	KDDCup99
sid:	271	did:	42769	C:	2	err:	35.9%	name:	Higgs
sid:	271	did:	43072	C:	2	err:	100.0%	name:	KDDCup09-Upselling